Outline

- Reinforcement Learning for Gomoku
- RL for News Recommendation
- RL for Text Generation
- Lab3 (In-class)
- Project3 – Blackjack (Homework)

Lab3 – Grid World

- Recall: What is Grid World problem?
  - States
  - Actions
  - Rewards
  - …

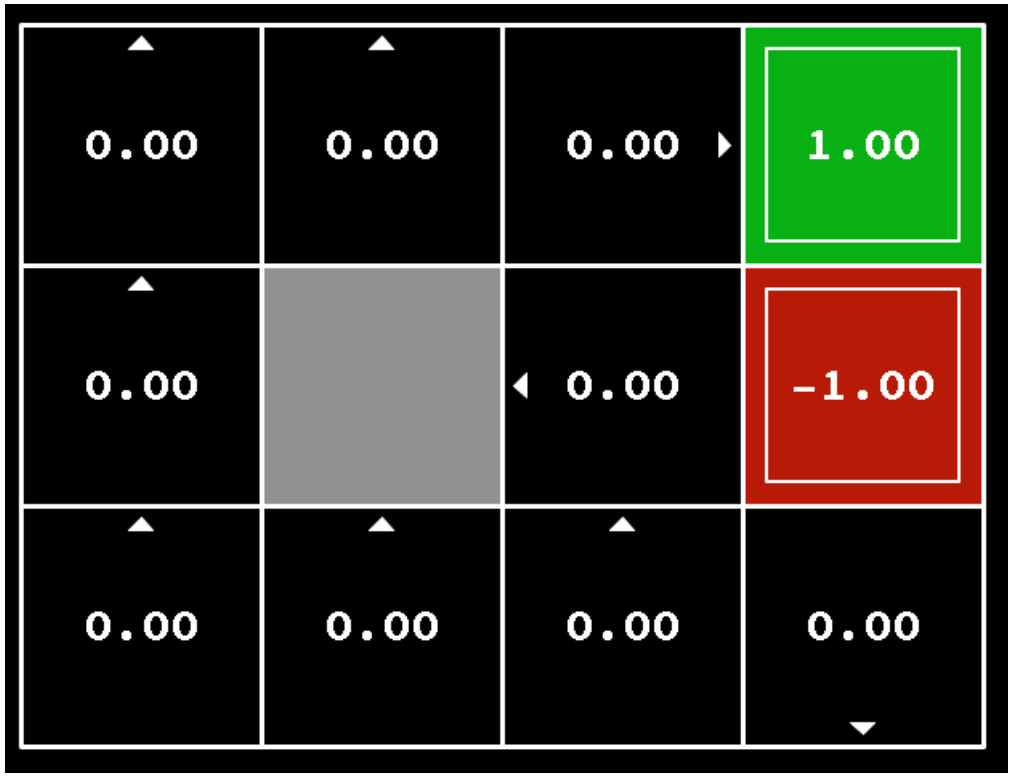

Lab3 – Grid World

- How about today’ s Grid World problem?
  - States
  - Actions:

    Up, Down, Left, Right
    Deterministically goes to next state
  - Rewards

0.0	A	0.0	B	0.0
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	B_{TO}	0.0
0.0	0.0	0.0	0.0	0.0
0.0	A_{TO}	0.0	0.0	0.0

$$R(s,a) = \begin{cases} 0.0, & taking\ a\ will\ stay\ in\ the\ grid\ world \\ -1.0, & taking\ a\ will\ jump\ out\ of\ the\ grid\ world \\ r, & current\ state\ is\ special \end{cases}$$

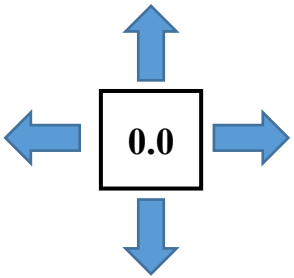

Lab3 – Grid World

- How about today’ s Grid World problem?
  - Value Iteration

$$V^*(s) = \max_{a \in A(s)} R(s,a) + \gamma \sum_{s'} P(s'|s,a) * V^*(s')$$

    - Synchronous Update
    - Asynchronous Update

0.0	A	0.0	B	0.0
0.0	0.0	0.0	0.0	0.0
0.0	0.0	0.0	B_{TO}	0.0
0.0	0.0	0.0	0.0	0.0
0.0	A_{TO}	0.0	0.0	0.0

Make sure you use synchronous update in this Lab.

Lab3 – Grid World

- How about today’ s Grid World problem?
  - Policy Evaluation

$$V^{\pi}(s) = R(s, \pi(s)) + \gamma \sum_{s'} P(s'|s, \pi(s)) * V^{\pi}(s')$$

  - Policy Improvement

$$\pi(s) = \arg \max_{a \in A(s)} Q(s, a)$$

0.0	A	**0.0**	B	**0.0**
0.0	**0.0**	**0.0**	**0.0**	**0.0**
0.0	**0.0**	**0.0**	B_{TO}	**0.0**
0.0	**0.0**	**0.0**	**0.0**	**0.0**
0.0	A_{TO}	**0.0**	**0.0**	**0.0**

Lab3

- Problem:
  - Solve the Grid World Problem based on MDP
- Requirement:
  - Print the iteration numbers and optimal values of all states using value iteration and policy iteration
- Address: http://10.192.9.85/contest/5/problem/01

Lab3

- Value Iteration

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp, an MDP with states S, actions $A(s)$, transition model $P(s'|s,a)$, rewards $R(s)$, discount γ .
ϵ, the maximum error allowed in the utility of all states

local variables: U, U', vectors of utilities for states in S, initially zero
δ, the maximum change in the utility of any stage in an iteration

repeat
  $U \leftarrow U'$; $\delta \leftarrow 0$
  for each state s **in** S **do**
    $U'[s] \leftarrow \max_{a \in A(s)} R(s,a) + \gamma \sum_{s'} P(s'|s,a)\, U[s']$
    $\delta \leftarrow \delta + |U'[s] - U[s]|$
until $\delta < \epsilon$
return U

Lab3

- Policy Iteration

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp, an MDP with states S, actions $A(s)$, transition model $P(s'|s,a)$
local variables: U, a vector of utilities for states in S, initially zero
π, a policy vector indexed by state, initially random

repeat
$U \leftarrow$ POLICY_EVALUATION(π, U, mdp)
$unchanged? \leftarrow$ true
for each state s **in** S **do**
if $\max_{a \in A(s)} Q(s,a) > Q(s,\pi[s])$ **then do**
$\pi[s] \leftarrow \arg\max_{a \in A(s)} Q(s,a)$
$unchanged? \leftarrow$ false
until $unchanged?$
return π

Think by yourself :
How to compute the
Q-value in this problem?

Outline

- Reinforcement Learning for Gomoku
- RL for News Recommendation
- RL for Text Generation
- Lab3 (In-class)
- Project3 – Blackjack (Homework)

Blackjack

- You need to submit your own version of code.
- You are encouraged to discuss with your group members.

  It might take some time to get familiar with all the supportive codes.
- Homework 3 is due on **23:55 pm, 09 Dec, 2020**